

**Rouibah Hanine**

**2025-2026**

# Report on Pneumonia vs. Normal Classification

## 1. Part 1 : Training a CNN Model Directly with Images

### 1.1. Data Preparation


**Dataset :** Images from the two classes (**Pneumonia and Normal**):

- **Training :** 3200 images for the **Pneumonia** class and 708 images for the **Normal** class
- **Validation :** 651 images for the **Pneumonia** and 621 images for the **Normal** class

**Class imbalance issue :** The Pneumonia and Normal classes were initially imbalanced.

Data augmentation techniques were applied to artificially generate more examples in both classes to balance them.

### Before Data Augmentation (training data only)




### Data Augmentation :


- **Rotation ±15°**
- **Vertical shift 5%**
- **Horizontal shift 5%**
- **Zoom 10%**
- **Horizontal flip**

We augmented 20% of the data for the Pneumonia class and 100% for the Normal class.

### After Data Augmentation



## Visualization of augmented data




**Image Resizing** : All images (original and generated) were resized to **150×150** pixels and saved for future use.

**Data Normalization** : Pixel values were normalized between **[0, 1]**.


## 1.2. CNN Architecture

- **Layers :**
  - **Convolutional Layers (Conv2D)** : Apply filters to the input image to extract spatial features such as edges, textures, and patterns.
  - **MaxPooling Layers** : Reduce the spatial dimensions of the feature maps, retaining the most important information while reducing computational complexity.
  - **Batch Normalization** : Normalizes activations to stabilize and accelerate training.
  - **Dropout** : Randomly deactivates a fraction of input units during training to prevent overfitting.
  - **Flatten** : Converts **2D** feature maps into a **1D** vector for input into fully connected layers.
  - **Fully Connected Dense Layers with Activation Functions** : **ReLU** is used in hidden layers to introduce non-linearity, while **Sigmoid** is used in the output layer for binary classification.
- **Training Parameters :**
  - **Epochs : 15**
  - **Optimizer : RMSprop (Root Mean Square Propagation)** was used to adjust the network weights, due to its ability to maintain an adaptive and stable learning rate.
  - **Learning Rate : 0.001**
  - **Early Stopping** : Monitors validation accuracy and stops training if no improvement is observed after **3 epochs**.
  - **Reduction of learning rate (ReduceLROnPlateau)** : Monitors validation accuracy and reduces the **learning rate** by a factor of **0.3** if no improvement is seen after **2 epochs**.


### 1.3. Training Results



### The Confusion Matrix :



### Receiver operating characteristic (ROC Curve):



- **Accuracy : 95.36 %**
- **Training Time : 2032.55 seconds**

## 2. Part 2 : Training a Machine Learning Model with Feature Extraction


### 2.1. Data Preparation

- **Dataset :** The same images used in training the **CNN** model were used here and for the next part, to ensure a consistent basis for comparison.
- **Extracted Features :**
  - **Pyramidal Histogram of Oriented Gradients (PHOG) :** Encodes spatial and orientation information of gradients, capturing the shape and structure of objects.
  - **Gabor Filters :** Capture texture features by analyzing frequency and orientation of the image at multiple scales.
  - **Discrete Cosine Transform (DCT) :** Transforms spatial data into the frequency domain, highlighting low-frequency components often critical for classification.
  - **Fourier Transform :** Converts spatial information into frequency components, useful for analyzing repetitive patterns or periodic structures.
- **Feature Vector**
  - Concatenation of the four methods : **[PHOG, Gabor, Fourier, DCT]**
  - This combination captures complementary features such as texture, frequency, and shape, allowing a comprehensive representation of the images.
- **Data Normalization :** Feature values were normalized between **[0, 1]** using the **Min-Max** method.


### 2.2. Model

- **Algorithm :** **SGDClassifier** with **log\_loss** loss function (logistic regression).
- **Training Parameters :**
  - **Epochs : 1000**
  - **Learning Rate : 0.0001** adaptive, which decreases based on model performance.
  - **Alpha = 0.0001** : Regularization strength to prevent overfitting.
  - **Penalty = 'elasticnet'** : Elastic Net regularization combining L1 (lasso) and L2 (ridge) penalties.
  - **L1\_ratio = 0.5 : 50 % L1 and 50 % L2** to balance the two types of regularization.
  - **Shuffle** : Shuffles the data at each iteration to improve convergence.
  - **Early stopping** : Monitors validation accuracy and stops training if no improvement is observed after **20 epochs**.
- Why Use **SGDClassifier** ?
  - **Performance on large Datasets** : **SGDClassifier** is better suited than **LogisticRegression** for large datasets, as it uses an incremental approach to adjust weights, making it faster and less memory-intensive.
  - **Control Over Regularization** : **SGDClassifier** allows direct specification of regularization parameters such as **I1** or **I2**, which is useful for preventing overfitting.


### 2.3. Training Results :



### The Confusion Matrix :



### Receiver operating characteristic (ROC Curve):



- **Accuracy : 93.63 %**
- **Training Time : 629.05 seconds**

### 3. Part 3 : Sequential Feature Processing for Machine Learning Models

#### 3.1. Data Preparation

- Feature Extraction Pipeline :

- Sequence : Gabor → Fourier → DCT → PHOG


- Reasoning :

- **Gabor Filters** : Extract initial texture and orientation features as a base.
- **Fourier Transform** : Analyzes periodic structures derived from **Gabor** features.
- **DCT** : Focuses on low-frequency components of the data transformed by **Fourier**, compressing relevant information.
- **PHOG** : Encodes spatial and gradient information from the processed data, providing shape-based features to complete the analysis.


#### 3.2. Model

- Algorithm : Identical to Part 2 (**SGDClassifier** with **log\_loss**).
- Model Parameters : The same parameters used in the previous part.


#### 3.3. Training Results :



#### The Confusion Matrix :



## Receiver operating characteristic (ROC Curve):



- **Accuracy : 86.87 %**
  - **Training Time : 72.62 seconds**
- 

## 4. Comparisons

### 4.1. Machine Learning Models

#### 1. Concatenation-Based Features

- **Accuracy : 93.63 %**
- **Execution Time : 629.05 seconds**

#### 2. Sequential Features :

- **Accuracy : 86.87 %**
- **Execution Time : 72.62 seconds**

**Observation :** The concatenation-based feature approach outperformed the sequential feature processing in terms of accuracy, at the cost of increased computation time.

### 4.2. Machine Learning vs. CNN

- **CNN Model**
  - **Accuracy : 95.36 %**
  - **Execution Time : 2032.55 seconds**
- **Best ML Model (Concatenation-Based Features)**
  - **Accuracy : 93.63 %**
  - **Execution Time : 629.05 seconds**

**Observation :** The **CNN** model provided the highest accuracy, demonstrating its superior ability to learn directly from images. However, it required significantly more computation time compared to the machine learning models.


---

## 5. Presentation of the Prediction Interface

In addition to the developed models, a user interface was implemented to enable real-time prediction.

This interface allows:

- The user to choose the model : **CNN** or **ML**.
- Upload an image.
- Directly display the classification result : **Sick** or **Not Sick**, accompanied by a confidence percentage.
- Save the analyzed image in dedicated folders, classified according to the prediction result.



## 6. Conclusion

The **CNN** model is highly effective for binary classification (**Pneumonia vs. Normal**), achieving the best accuracy (**95.36%**).

Among the **machine learning models**, the concatenation-based feature approach (**PHOG, Gabor, Fourier, DCT**) obtained better results than the **sequential approach**, with a notable accuracy of **93.63%**.

For scenarios where **execution time** is a critical factor, the sequential feature processing model (**86.87%** accuracy in **72.62 seconds**) can be a **viable alternative**.

**END**